



Introduction to Data Mining

Web Chapter Exploring Data

by Michael Hahsler

Based in Slides by Tan, Steinbach,
Karpatne, Kumar



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).

R Code Examples

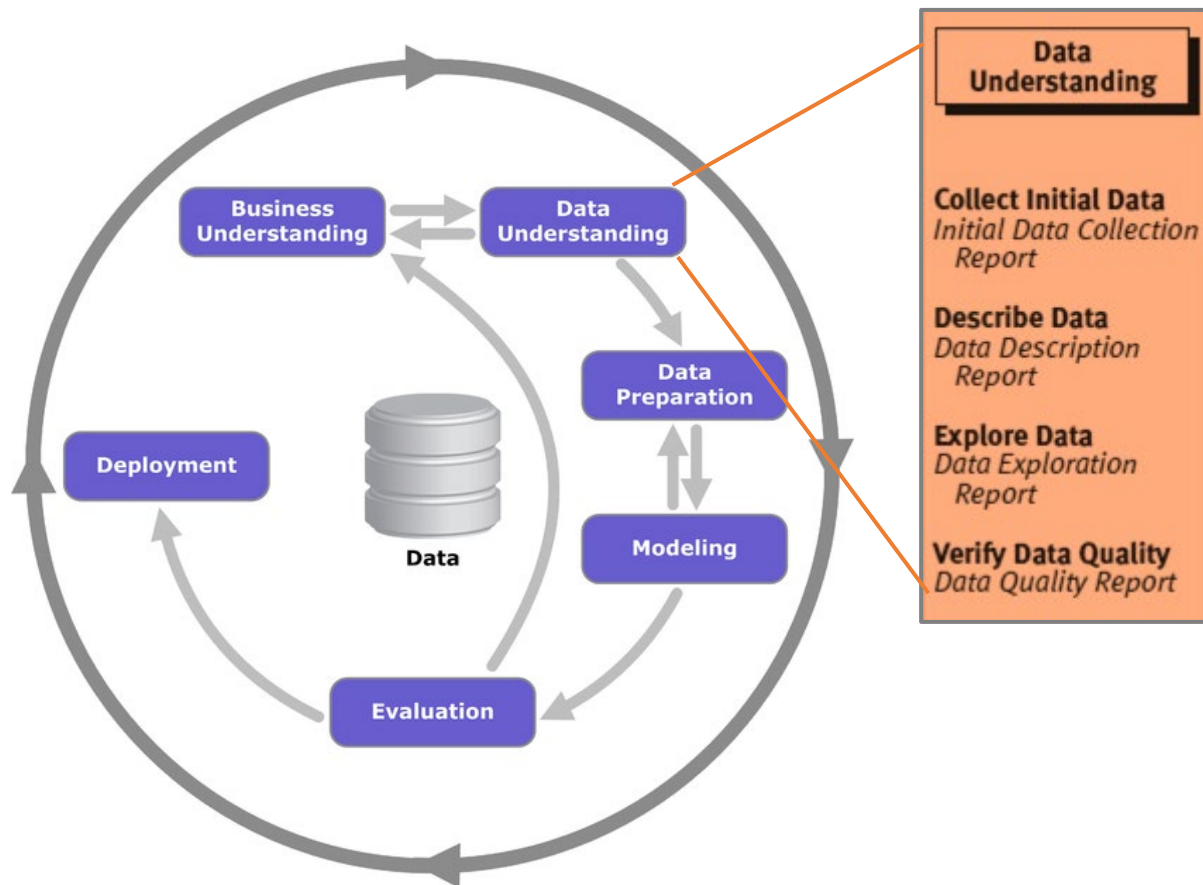
- Available R Code examples are indicated on slides by the R logo



- The Examples are available at [https://mhahsler.github.io/Introduction to Data Mining R Examples/](https://mhahsler.github.io/Introduction%20to%20Data%20Mining%20R%20Examples/)



Exploring Data in the Data Mining Process



Topics

- **Exploratory Data Analysis**
- Summary Statistics
- Visualization



What is Data Exploration?

“A preliminary exploration of the data to better understand its characteristics.”

- Key motivations of data exploration include
 - Helping to select the right tool for preprocessing or analysis
 - Making use of humans' abilities to recognize patterns.
 - People can recognize patterns not captured by data analysis tools
- Related to the area of Exploratory Data Analysis (EDA)
 - Created by statistician John Tukey
 - Seminal book is "Exploratory Data Analysis" by Tukey
 - A nice online introduction can be found in Chapter 1 of the NIST Engineering Statistics Handbook
 - <http://www.itl.nist.gov/div898/handbook/index.htm>

Topics

- Exploratory Data Analysis
- **Summary Statistics**
- Visualization



Summary Statistics



Summary statistics are numbers that summarize properties of the data



Summarized properties include location and spread for continuous data

Examples: location - mean
 spread - standard deviation



Most summary statistics can be calculated in a single pass through the data

Frequency and Mode

- The frequency of an attribute value is the percentage of time the value occurs in the data set
 - For example, given the attribute 'gender' and a representative population of people, the gender 'female' occurs about 50% of the time.
- The mode of an attribute is the most frequent attribute value
- The notions of frequency and mode are typically used with **categorical data**.



Measures of Location: Mean and Median

- For quantitative features.
- The mean is the most common measure of the location of a set of points.
- However, the mean is very sensitive to outliers.
- Thus, the median or a trimmed mean is also commonly used.

$$\text{mean}(x) = \bar{x} = \frac{1}{m} \sum_{i=1}^m x_i$$

$$\text{median}(x) = \begin{cases} x_{(r+1)} & \text{if } m \text{ is odd, i.e., } m = 2r + 1 \\ \frac{1}{2}(x_{(r)} + x_{(r+1)}) & \text{if } m \text{ is even, i.e., } m = 2r \end{cases}$$

Measures of Spread: Range and Variance

- Range is the difference between the max and min
- The variance or standard deviation is the most common measure of the spread of a set of points.

$$\text{variance}(x) = s_x^2 = \frac{1}{m-1} \sum_{i=1}^m (x_i - \bar{x})^2$$

- However, this is also sensitive to outliers, so that other measures are often used.

$$\text{AAD}(x) = \frac{1}{m} \sum_{i=1}^m |x_i - \bar{x}|$$

$$\text{MAD}(x) = \text{median}\left(\{|x_1 - \bar{x}|, \dots, |x_m - \bar{x}|\}\right)$$

$$\text{interquartile range}(x) = x_{75\%} - x_{25\%}$$




Percentiles

- Given an ordinal or continuous attribute x and a number p between 0 and 100, the p^{th} percentile is a value $x_{p\%}$ of x such that $p\%$ of the observed values of x are less than $x_{p\%}$.
- Example: the 50th percentile is the value $x_{50\%}$ such that 50% of all values of x are less than $x_{50\%}$.

Percentiles

Median – 50% of the cases has a smaller value & 50% are larger



Multivariate Summary Statistics

Object	x_1	x_2
1	12	15
2	2	4
...
m	18	4

- Covariance between features i and j

$$s_{ij} = \frac{1}{m-1} \sum_{k=1}^m (x_{ki} - \bar{x}_i)(x_{kj} - \bar{x}_j)$$

- Correlation

$$r_{ij} = \frac{s_{ij}}{s_i s_j}$$

- s_i is the variance of feature i

Correlation

- Correlation measures the (linear) relationship between two variables.
- To compute Pearson correlation (Pearson's Product Moment Correlation), we standardize data objects, p and q , and then take their dot product

$$\rho = \frac{\text{cov}(X, Y)}{\text{sd}(X) \text{sd}(Y)}$$


- Estimation:

$$r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}}$$

- Correlation is often used as a measure of similarity.

Visually Evaluating Correlation

Scatter plots showing the
similarity from -1 to 1 .



Rank Correlation

- Measure the degree of similarity between two ratings (e.g., ordinal data).
- Is more robust against outliers and does not assume normality of data or linear relationship like Pearson Correlation.
- Measures (all are between -1 and 1)
 - Spearman's Rho: Pearson correlation between ranked variables.
 - Kendall's Tau

$$\tau = \frac{N_s - N_d}{\frac{1}{2}n(n-1)}$$

N_s ... concordant pair
 N_d ... discordant pair

- Goodman and Kruskal's Gamma

$$\gamma = \frac{N_s - N_d}{N_s + N_d}$$





Topics

- Exploratory Data Analysis
- Summary Statistics
- **Visualization**



Visualization



Visualization is the conversion of data into a visual or tabular format so that the characteristics of the data and the **relationships among data items or attributes** can be analyzed or reported.



Visualization of data is one of the most powerful and appealing techniques for data exploration.


Humans have a well-developed ability to analyze large amounts of information that is presented visually

Can detect general patterns and trends

Can detect outliers and unusual patterns

Example: Sea Surface Temperature

- The following shows the Sea Surface Temperature (SST) for July 1982
- Tens of thousands of data points are summarized in a single figure



Representation

- Is the mapping of information to a visual format
- Data objects, their attributes, and the relationships among data objects are translated into graphical elements such as points, lines, shapes, and colors.
- Example:
 - Objects are often represented as points
 - Their attribute values can be represented as the position of the points or the characteristics of the points, e.g., color, size, and shape
 - If position is used, then the relationships of points, i.e., whether they form groups or a point is an outlier, is easily perceived.

Arrangement

- Is the placement of visual elements within a display
- Can make a large difference in how easy it is to understand the data

Example:

	1	2	3	4	5	6
1	0	1	0	1	1	0
2	1	0	1	0	0	1
3	0	1	0	1	1	0
4	1	0	1	0	0	1
5	0	1	0	1	1	0
6	1	0	1	0	0	1
7	0	1	0	1	1	0
8	1	0	1	0	0	1
9	0	1	0	1	1	0

	6	1	3	2	5	4
4	1	1	1	0	0	0
2	1	1	1	0	0	0
6	1	1	1	0	0	0
8	1	1	1	0	0	0
5	0	0	0	1	1	1
3	0	0	0	1	1	1
9	0	0	0	1	1	1
1	0	0	0	1	1	1
7	0	0	0	1	1	1

The Iris Dataset

Many of the exploratory data techniques are illustrated with the Iris Plant data set.

- Included as a demo dataset in many tools (R, scikit-learn, Rapidminer, ...).
- Can be obtained from the UCI Machine Learning Repository <http://www.ics.uci.edu/~mlearn/MLRepository.html>
- From the statistician R.A. Fisher
- 150 flowers, three types (classes).
- Four (non-class) attributes

	▲ Sepal.Length ▾	Sepal.Width ▾	Petal.Length ▾	Petal.Width ▾	Species ▾
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa

...

Iris Versicolor



Iris Virginica





Iris Setosa



Distribution: Histograms


- Usually shows the distribution of values of a single variable
- Divide the values into bins and show a bar plot of the number of objects in each bin.
- The height of each bar indicates the number of objects
- Shape of histogram depends on the number of bins

Example: Petal Width (10 and 20 bins, respectively)




Distribution Box Plots

- Invented by J. Tukey
- Simplified version of a PDF/histogram.





PDF or histogram



Examples of Box Plots

- Box plots can be used to compare attributes or subgroups.


	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa



Two-Dimensional Histograms


- Show the joint distribution of the values of two attributes

Example: petal width and petal length. What does this tell us?




Scatter Plots

- Attributes values determine the position
- Two-dimensional scatter plots most common, but can have three-dimensional scatter plots
- Often additional attributes can be displayed by using the size, shape, and color of the markers that represent the objects




Scatter Plot Array of Iris Attributes



Contour Plots

- Useful when a continuous attribute is measured on a **spatial grid**
- They partition the plane into regions of similar values
- The contour lines that form the boundaries of these regions connect points with equal values
- The most common example is contour maps of elevation
- Can also display temperature, rainfall, air pressure, etc.




Matrix Plots


- Can plot a data matrix
- Can be useful when objects are sorted according to class
- Typically, the attributes are normalized to prevent one attribute from dominating the plot
- Plots of similarity or distance matrices can also be useful for visualizing the relationships between objects

Example: The Iris Data Matrix

Deviation from feature mean



Example: The Iris Correlation Matrix





Parallel Coordinates



- Used to plot the attribute values of high-dimensional data
- Instead of using perpendicular axes, use a set of parallel axes
- The attribute values of each object are plotted as a point on each corresponding coordinate axis and the points are connected by a line
- Thus, each object is represented as a line
- Often, the lines representing a distinct class of objects group together, at least for some attributes
- Ordering of attributes is important in seeing such groupings

Example: Parallel Coordinates Plots for Iris Data

Reordered features



Other Visualization Techniques

- Star Plots
 - Similar approach to parallel coordinates, but axes radiate from a central point
 - The line connecting the values of an object is a polygon
- Chernoff Faces
 - Approach created by Herman Chernoff
 - This approach associates each attribute with a characteristic of a face
 - The values of each attribute determine the appearance of the corresponding facial characteristic
 - Each object becomes a separate face
 - Relies on human's ability to distinguish faces

Example: Star Plots for Iris Data

- Setosa



1



2



3




4




5


- Versicolor




51




52



53




54




55


- Virginica




101




102



103



104



105

Example: Chernoff Faces for Iris Data

- Setosa



1



2



3



4



5

- Versicolor



51



52



53



54



55

- Virginica



101



102



103



104



105



Conclusion

- Exploring data is the first step when working with data.
- The goal is to:
 1. Understand what data is available.
 2. Assess data distributions and how variables relate to each other.
 3. Assess data quality.
- Understanding the data is necessary to decide on data preparation and modeling.