



# Introduction to Data Mining

## Chapter 4 Classification – Alternative Techniques

---

by Michael Hahsler

Based in Slides by Tan,  
Steinbach, Karpatne, Kumar



# R Code Examples

- Available R Code examples are indicated on slides by the R logo



- The Examples are available at [https://mhahsler.github.io/Introduction to Data Mining R Examples/](https://mhahsler.github.io/Introduction%20to%20Data%20Mining%20R%20Examples/)





# Topics

---

- Other Classification Methods
  - **Rule-Based Classifier**
  - Nearest Neighbor Classifier
  - Naive Bayes Classifier
  - Artificial Neural Networks
  - Support Vector Machines
  - Ensemble Methods
- Class Imbalance Problem

# Rule-Based Classifier

- Classify records by using a collection of “if...then...” rules
- Rule:  $(Condition) \rightarrow y$ 
  - Condition is a conjunctions of attributes called LHS, antecedent or condition
  - $y$  is the class label called RHS or consequent
- Examples of classification rules for an animal dataset:
  - $(Blood\ Type = Warm) \wedge (Lay\ Eggs = Yes) \rightarrow Birds$
  - $(Taxable\ Income < 50K) \wedge (Refund = Yes) \rightarrow Evade = No$

# Using a Rule-Based Classifier

A rule  $R$  **covers** an instance  $x$  if the attributes of the instance satisfy the condition of the rule. Such a rule can be used for classification.

R1: (Give Birth = no)  $\wedge$  (Can Fly = yes)  $\rightarrow$  Birds

R2: (Give Birth = no)  $\wedge$  (Live in Water = yes)  $\rightarrow$  Fishes

R3: (Give Birth = yes)  $\wedge$  (Blood Type = warm)  $\rightarrow$  Mammals

R4: (Give Birth = no)  $\wedge$  (Can Fly = no)  $\rightarrow$  Reptiles

R5: (Live in Water = sometimes)  $\rightarrow$  Amphibians

Rule base

Name	Blood Type	Give Birth	Can Fly	Live in Water	Class
hawk	warm	no	yes	no	?
grizzly bear	warm	yes	no	no	?

The rule R1 covers: *hawk*  $\rightarrow$  *Bird*

The rule R3 covers: *grizzly bear*  $\rightarrow$  *Mammal*

# Ordered Rule Set vs. Voting

- Rules are rank ordered according to their priority
  - An ordered rule set is known as a decision list
- When a test record is presented to the classifier
  - It is assigned to the class label of the highest ranked rule it has triggered (R3 is selected below -> Amphibians)
  - If none of the rules fired, it is assigned to the default class

R1: (Give Birth = no)  $\wedge$  (Can Fly = yes)  $\rightarrow$  Birds

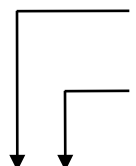
R2: (Give Birth = yes)  $\wedge$  (Blood Type = warm)  $\rightarrow$  Mammals

R3: (Live in Water = sometimes)  $\rightarrow$  Amphibians

R4: (Give Birth = no)  $\wedge$  (Can Fly = no)  $\rightarrow$  Reptiles

R5: (Give Birth = no)  $\rightarrow$  Amphibians

Rule base

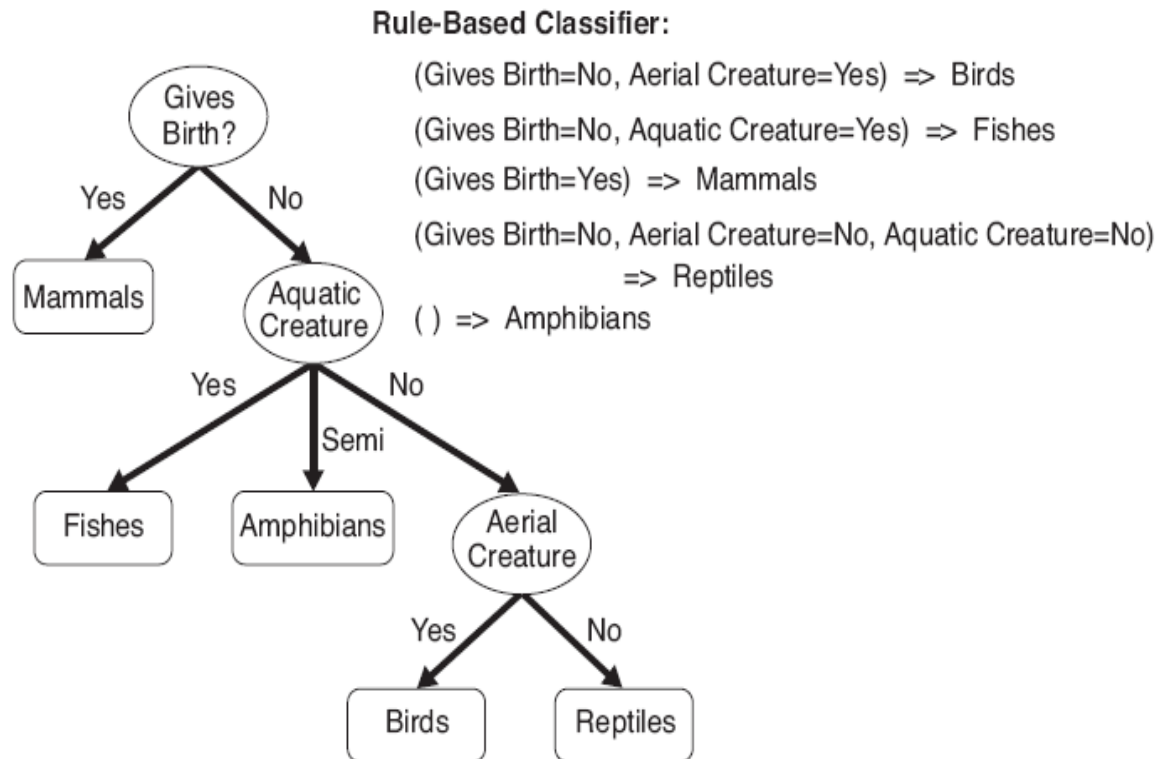


Name	Blood Type	Give Birth	Can Fly	Live in Water	Class
turtle	cold	no	no	sometimes	?

- Alternative: (weighted) voting by all matching rules (-> Amphibians)

R3, 4 and 5  
cover the  
observation

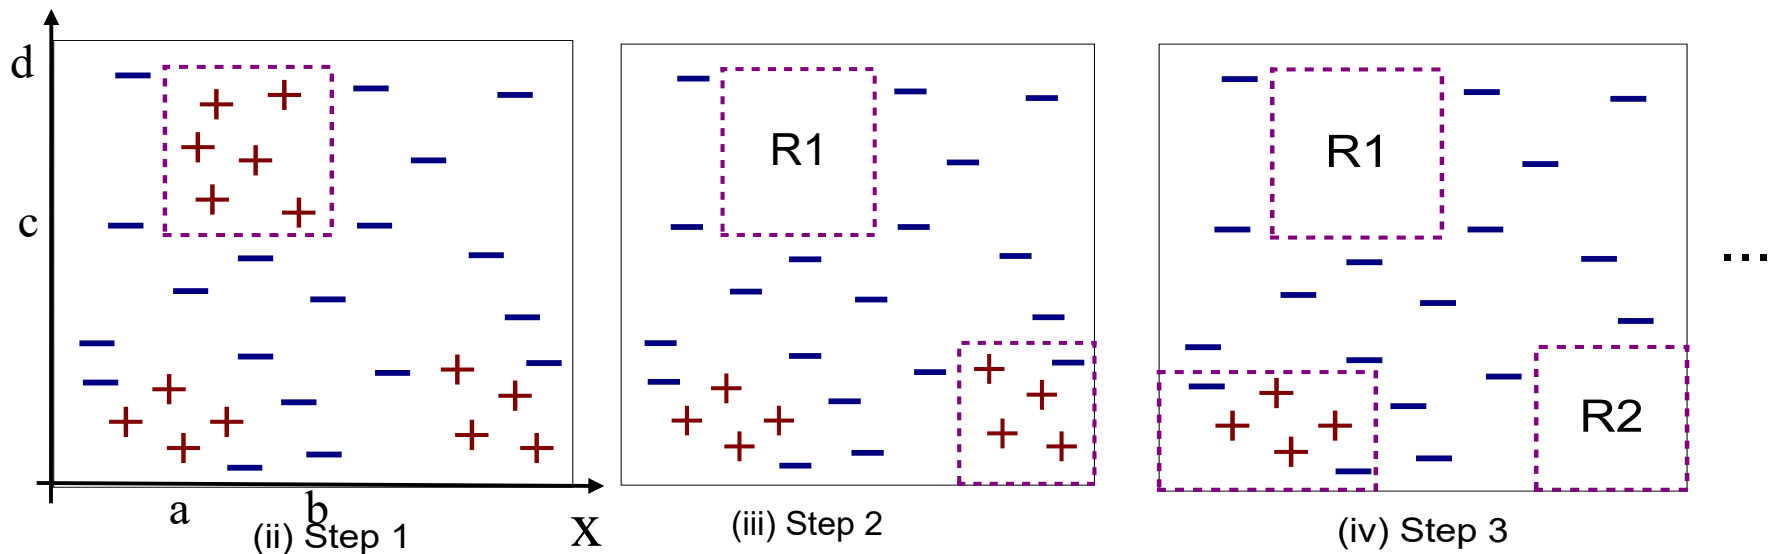
# Rules From Decision Trees



- Rules are created by reading the decisions in tree branches from the root to a final node.
- Rule set contains as much information as the tree.
- Rules can be simplified (similar to pruning of the tree).
- Example: C4.5rules

# Direct Methods of Rule Generation

- Extract rules directly from the data
- Sequential Covering (Example: try to cover class +)



$$R1: a > x > b \wedge c > y > d \rightarrow \text{class } +$$



# Advantages of Rule-Based Classifiers

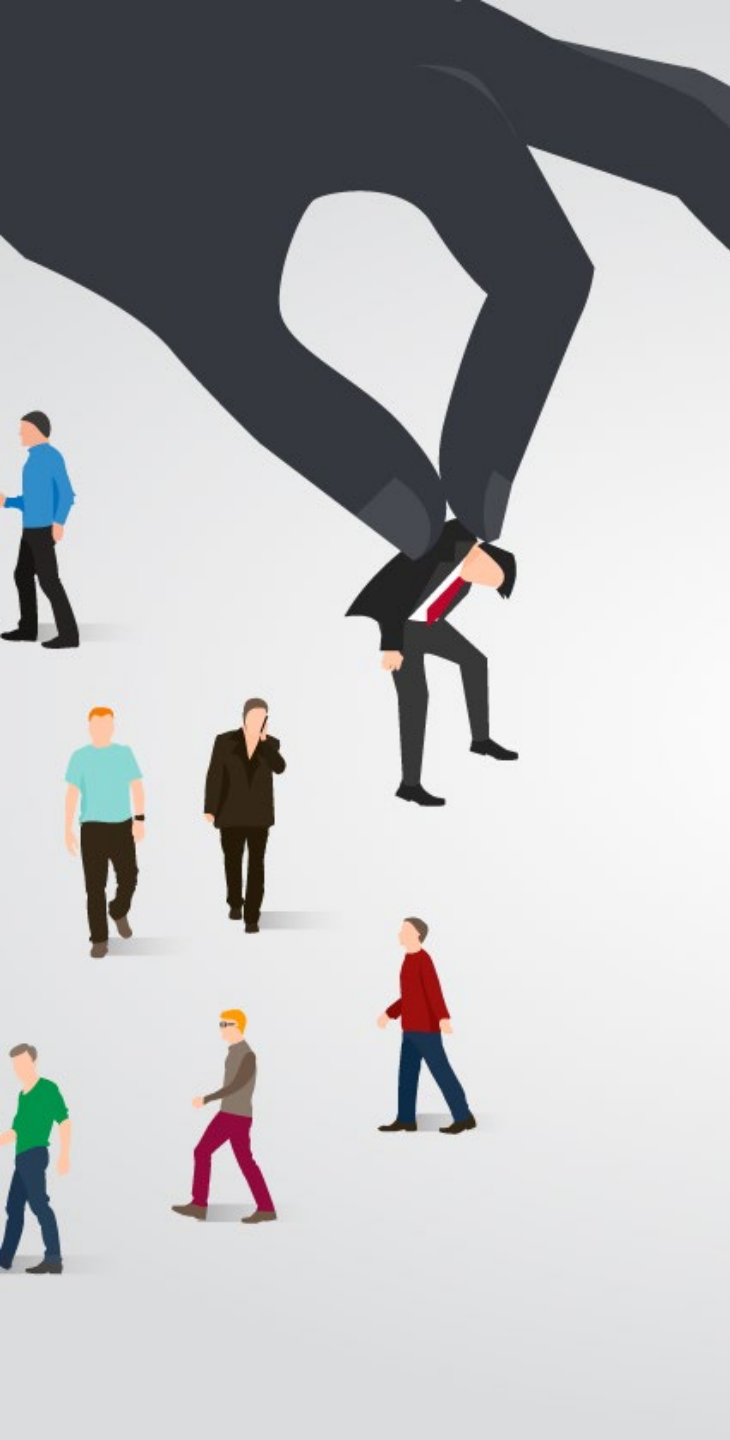
As expressive  
as decision  
trees

Easy to  
interpret

Easy to  
generate

Can classify  
new instances  
rapidly

Performance  
comparable to  
decision trees



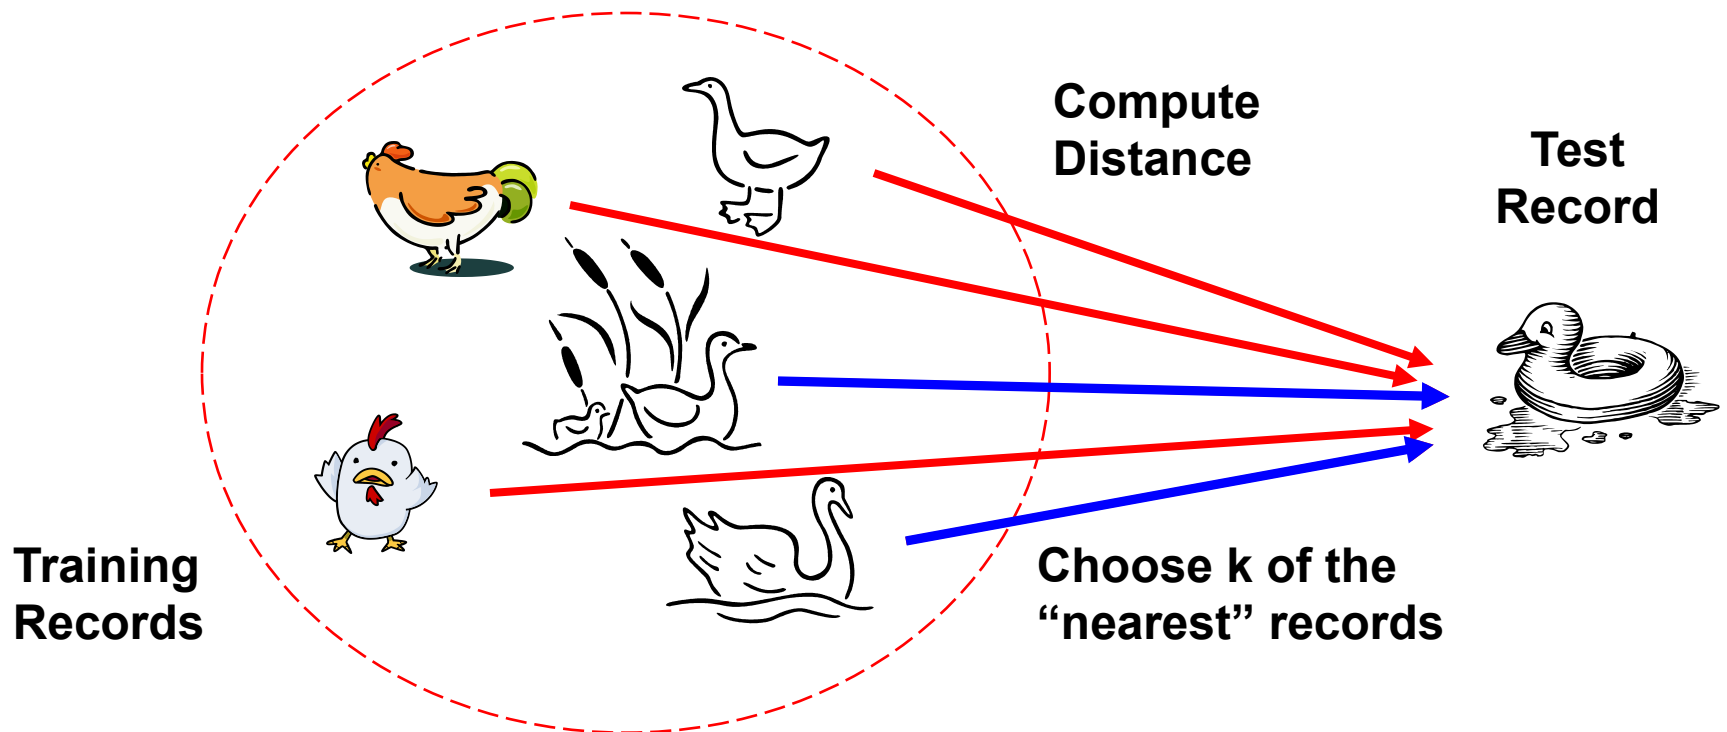
# Topics

- Rule-Based Classifier
- **Nearest Neighbor Classifier**
- Naive Bayes Classifier
- Artificial Neural Networks
- Support Vector Machines
- Ensemble Methods

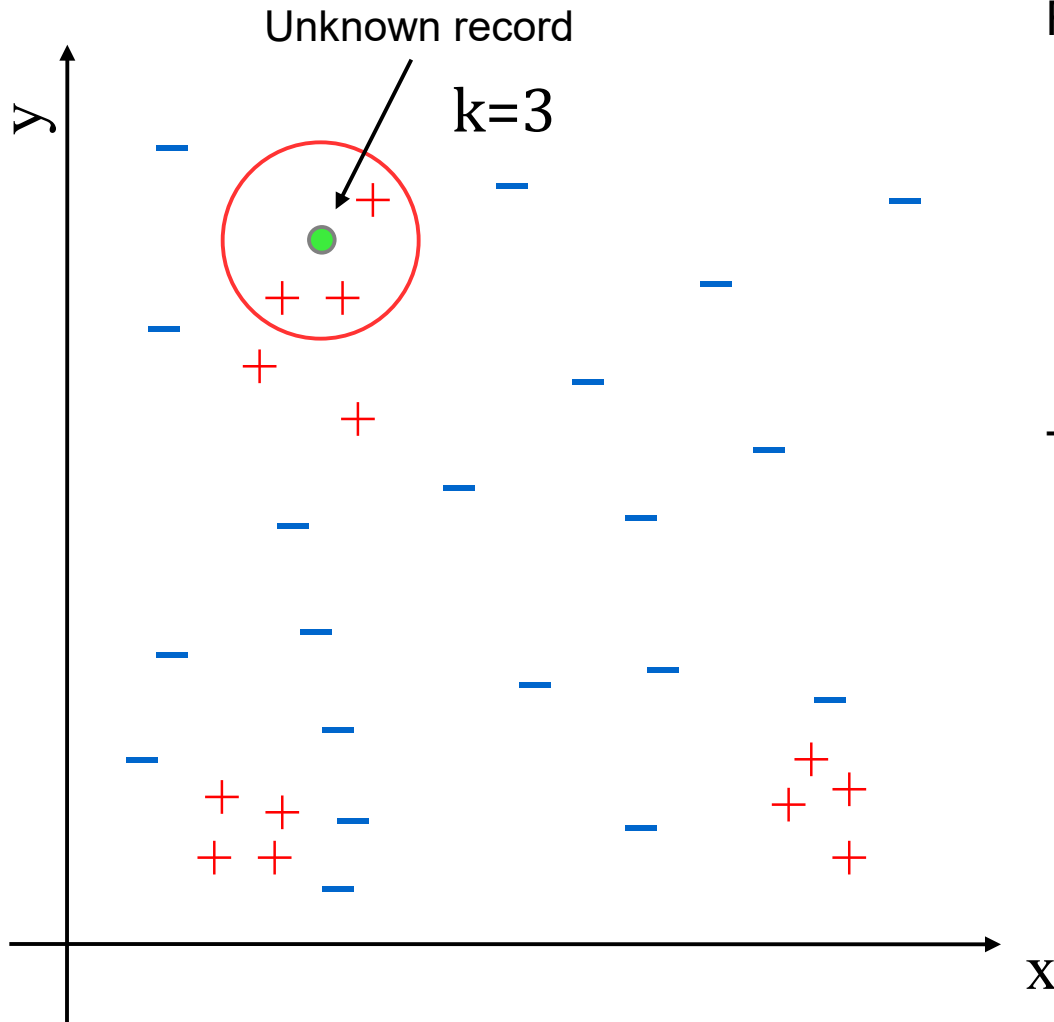
# Nearest Neighbor Classifiers

- Basic idea:

- If it walks like a duck, quacks like a duck, then it's probably a duck



# Nearest-Neighbor Classifiers




Requires three things

- The set of stored records
- Distance Metric to compute distance between records
- The value of  $k$ , the number of nearest neighbors to retrieve


To classify an unknown record:

- Compute distance to other training records
- Identify  $k$  nearest neighbors
- Use class labels of nearest neighbors to determine the class label of unknown record (e.g., by taking majority vote)


# Definition of Nearest Neighbor



(a) 1-nearest neighbor



(b) 2-nearest neighbor



(c) 3-nearest neighbor

K-nearest neighbors of a record  $x$  are data points that have the  $k$  smallest distance to  $x$

# Nearest Neighbor Classification

- Compute distance between two points:
  - Euclidean distance

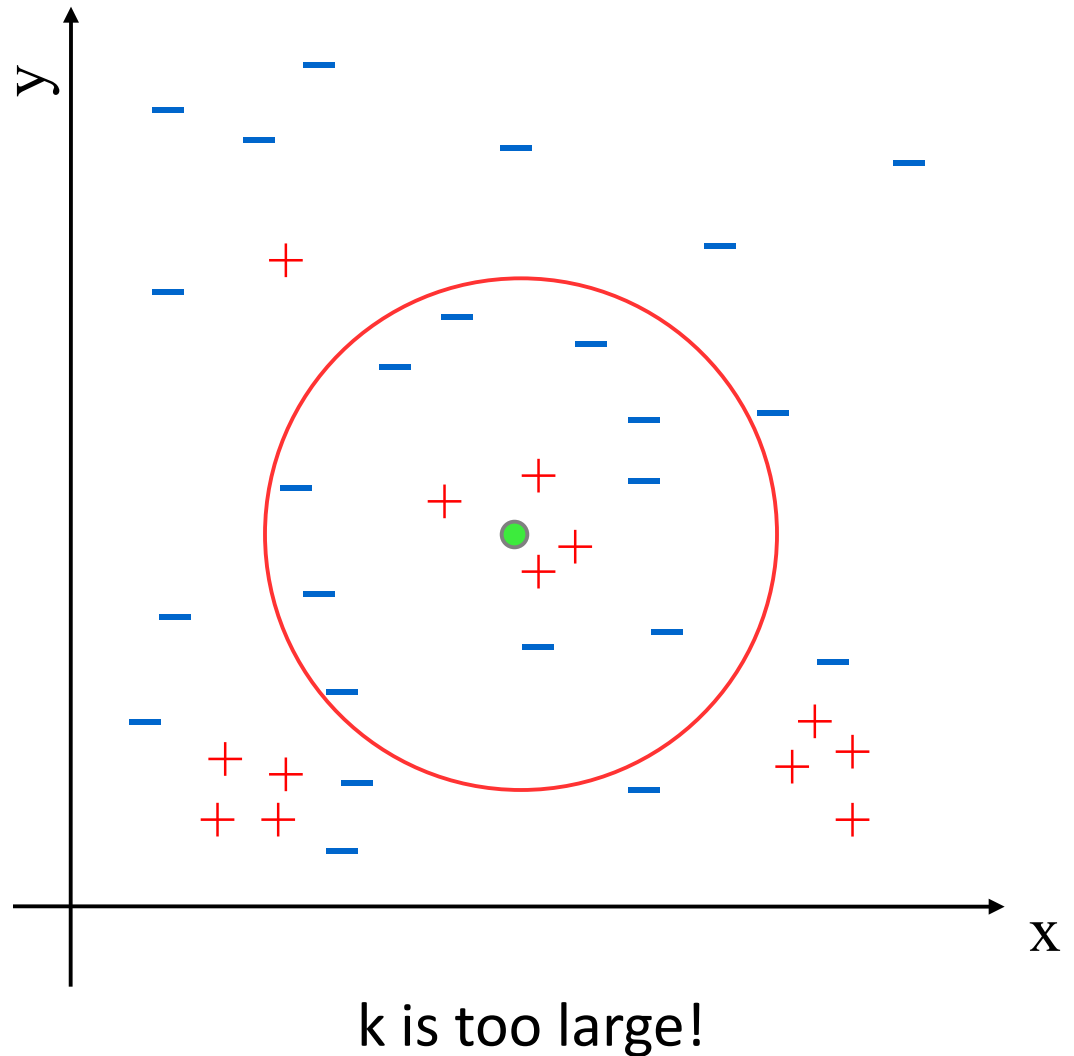
$$d(\mathbf{p}, \mathbf{q}) = \sqrt{\sum_i (p_i - q_i)^2}$$

- Note: This means that the data needs to be scaled!
- Determine the class from nearest neighbor list
  - take the majority vote of class labels among the k-nearest neighbors
  - Weigh the vote according to distance (e.g., weight factor  $w = 1/d^2$ )

# Nearest Neighbor Classification...

- Choosing the value of  $k$ :

- If  $k$  is too small, sensitive to noise points
- If  $k$  is too large, neighborhood may include points from other classes

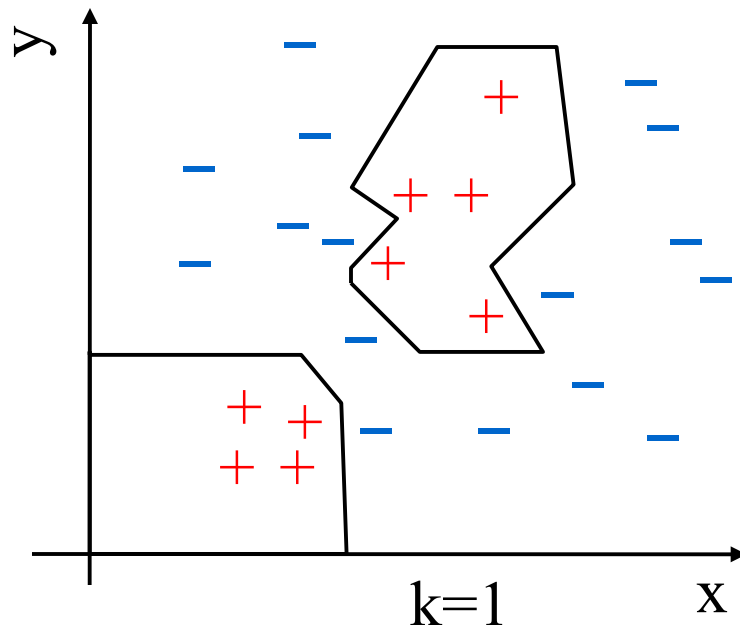


# Nearest neighbor Classification...

k-NN classifiers are lazy learners

- It does not build models explicitly (unlike eager learners such as decision trees)
- Needs to store all the training data
- Classifying unknown records are relatively expensive (find the k-nearest neighbors)

**Advantage:** Can create arbitrary non-linear decision boundaries.







# Topics

- Rule-Based Classifier
- Nearest Neighbor Classifier
- **Naive Bayes Classifier**
- Artificial Neural Networks
- Support Vector Machines
- Ensemble Methods

# Bayes' Rule

- The product rule gives us two ways to factor a joint distribution:

$$P(A, B) = P(A|B)P(B) = P(B|A)P(A)$$

- Therefore,

Posterior Prob.

Prior Prob.

$$P(A|B) = \frac{P(B|A) P(A)}{P(B)}$$

- Why is this useful?

- Can get diagnostic probability  $P(\text{cavity} \mid \text{toothache})$  from causal probability  $P(\text{toothache} \mid \text{cavity})$
- We can update our beliefs based on evidence.
- Important tool for probabilistic inference .

# Example of Bayes Theorem

- A doctor knows that meningitis causes stiff neck 50% of the time →  $P(S|M)=.5$
- Prior probability of any patient having meningitis is  $P(M) = 1/50,000=0.00002$
- Prior probability of any patient having stiff neck is  $P(S) = 1/20=0.05$
- If a patient has stiff neck, what's the probability he/she has meningitis?

$$P(M | S) = \frac{P(S | M) P(M)}{P(S)} = \frac{.5 \times 0.00002}{0.05} = 0.0002$$

Increases the probability by x10!

# Bayesian Classifiers

- Consider each attribute and class label as a random variable.
- Classification problem: Given a record with attributes  $(A_1, A_2, \dots, A_n)$  predict class  $C$ .
- This can be done by finding the most likely class that has the largest

$$P(C | A_1, A_2, \dots, A_n)$$

# Bayesian Classifiers

- Compute the posterior probability  $P(C | A_1, A_2, \dots, A_n)$  for all values of  $C$  using the Bayes theorem

$$\max_C P(C | A_1, A_2, \dots, A_n) = \max_C \frac{P(A_1, A_2, \dots, A_n | C) P(C)}{P(A_1, A_2, \dots, A_n)}$$

- Choose value of  $C$  that maximizes  $P(C | A_1, A_2, \dots, A_n)$

this is a constant!

- Equivalent to choosing value of  $C$  that maximizes  $\max_C P(A_1, A_2, \dots, A_n | C) P(C)$

- How do we estimate  $P(A_1, A_2, \dots, A_n | C)$ ?

# Naïve Bayes Classifier

Assume independence among attributes A when class is given:

$$P(A_1, A_2, \dots, A_n | C) = P(A_1 | C) P(A_2 | C) \dots P(A_n | C) = \prod_i P(A_i | C)$$

We can estimate  $P(A_i | C_j)$  for all  $A_i$  and  $C_j$ .

A new observation is classified to  $C_j$  such that:

$$\max_j P(C_j) \prod_i P(A_i | C_j)$$

# How to Estimate Probabilities from Data?

- Class:  $P(C_j) = N_{C_j} / N$

e.g.,  $P(C=\text{No}) = 7/10$ ,  
 $P(C=\text{Yes}) = 3/10$

- For discrete attributes:

$$P(A_i | C_j) = \frac{|A_{ij}|}{N_{C_j}}$$

where  $|A_{ij}|$  is number of instances having attribute  $A_i$  and belongs to class  $C_j$

e.g.

$P(\text{Status}=\text{Married} \mid C=\text{No}) = 4/7$

$P(\text{Refund}=\text{Yes} \mid C=\text{Yes})=0$

<i>Tid</i>	Refund	Marital Status	Taxable Income	Class
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

# How to Estimate Probabilities from Data?

For continuous attributes:

- Discretize the range into bins
  - one ordinal attribute per bin
  - violates independence assumption
- Two-way split:  $(A < v)$  or  $(A > v)$ 
  - choose only one of the two splits as new attribute
- Probability density estimation
  - Assume attribute follows a normal distribution
  - Use data to estimate parameters of distribution (e.g., mean and standard deviation)
  - Once probability distribution is known, can use it to estimate the conditional probability  $P(A_i | C_j)$



# Example of Naïve Bayes Classifier

**Given a Test Record what is the most likely class?**

$$X = (\text{Refund} = \text{No}, \text{Married}, \text{Income} = 120K)$$

naive Bayes Classifier:

$$P(\text{Refund}=\text{Yes}|\text{No}) = 3/7$$

$$P(\text{Refund}=\text{No}|\text{No}) = 4/7$$

$$P(\text{Refund}=\text{Yes}|\text{Yes}) = 0$$

$$P(\text{Refund}=\text{No}|\text{Yes}) = 1$$

$$P(\text{Marital Status}=\text{Single}|\text{No}) = 2/7$$

$$P(\text{Marital Status}=\text{Divorced}|\text{No}) = 1/7$$

$$P(\text{Marital Status}=\text{Married}|\text{No}) = 4/7$$

$$P(\text{Marital Status}=\text{Single}|\text{Yes}) = 2/7$$

$$P(\text{Marital Status}=\text{Divorced}|\text{Yes}) = 1/7$$

$$P(\text{Marital Status}=\text{Married}|\text{Yes}) = 0$$

For taxable income:

If class=No: sample mean=110

sample variance=2975

If class=Yes: sample mean=90

sample variance=25

$$\begin{aligned} P(X|\text{Class}=\text{No}) &= P(\text{Refund}=\text{No}|\text{Class}=\text{No}) \\ &\quad * P(\text{Married}|\text{Class}=\text{No}) \\ &\quad * P(\text{Income}=120K|\text{Class}=\text{No}) \\ &= 4/7 * 4/7 * 0.0072 = 0.0024 \end{aligned}$$

$$\begin{aligned} P(X|\text{Class}=\text{Yes}) &= P(\text{Refund}=\text{No}|\text{Class}=\text{Yes}) \\ &\quad * P(\text{Married}|\text{Class}=\text{Yes}) \\ &\quad * P(\text{Income}=120K|\text{Class}=\text{Yes}) \\ &= 1 * 0 * 1.2 * 10^{-9} = 0 \end{aligned}$$

Since  $P(X|\text{No})P(\text{No}) > P(X|\text{Yes})P(\text{Yes})$

Therefore  $P(\text{No}|X) > P(\text{Yes}|X)$

=> Class = No

# Naïve Bayes Classifier

Probability estimation:

Original:  $P(A_i | C_j) = \frac{N_{ij}}{N_j}$

Issue: If one of the conditional probabilities is zero, then the entire expression becomes zero.

Laplace:  $P(A_i | C_j) = \frac{N_{ij}+1}{N_j+c}$

c: number of classes

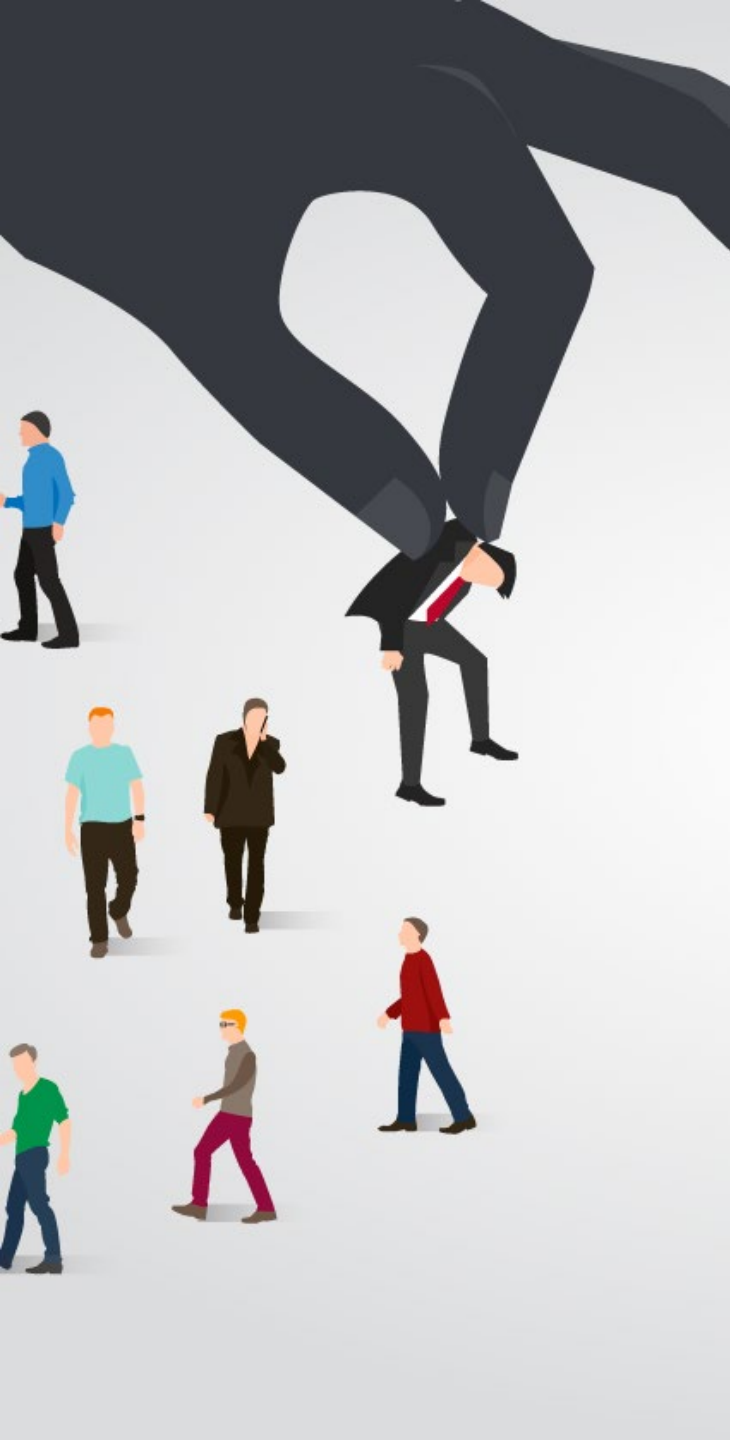
p: prior probability

m-estimate:  $P(A_i | C_j) = \frac{N_{ij}+mp}{N_j+m}$

m: parameter

# Naïve Bayes (Summary)


- Robust to isolated noise points
- Handle missing values by ignoring the instance during probability estimate calculations
- Robust to irrelevant attributes
- Independence assumption may not hold for some attributes
  - Use other techniques such as Bayesian Belief Networks (BBN)



# Topics

- Rule-Based Classifier
- Nearest Neighbor Classifier
- Naive Bayes Classifier
- **Artificial Neural Networks**
- Support Vector Machines
- Ensemble Methods

# General Structure of ANN



# Algorithm for learning ANN


- Initialize the weights ( $w_0, w_1, \dots, w_k$ )
- Adjust the weights in such a way that the output of ANN is consistent with class labels of training examples

—Objective function:

$$E = \sum_i [Y_i - f(w_i, X_i)]^2$$

- Find the weights  $w_i$ 's that minimize the above objective function.  
Methods: backpropagation algorithm, gradient descend

# Deep Learning / Deep Neural Networks



- Needs lots of data + computation (GPU)
- Applications: computer vision, speech recognition, natural language processing, audio recognition, machine translation, bioinformatics, ...
- Tools: Keras, Tensorflow and many others.
- Related: Deep belief networks, recurrent neural networks (RNN), convolutional neural network (CNN)




# Topics

- Rule-Based Classifier
- Nearest Neighbor Classifier
- Naive Bayes Classifier
- Artificial Neural Networks
- **Support Vector Machines**
- Ensemble Methods




# Support Vector Machines




Find a linear hyperplane (decision boundary) that will separate the data

# Support Vector Machines




One Possible Solution

# Support Vector Machines




Another possible solution

# Support Vector Machines



Other possible solutions


# Support Vector Machines



Which one is better? B1 or B2?

How do you define better?

# Support Vector Machines




Find hyperplane **maximizes** the margin => B1 is better than B2

Larger margin = more robust = less expected generalization error

# Support Vector Machines


What if the problem is not linearly separable?



- Use slack variables to account for violations
- Use hyperplane that minimizes slack


# Nonlinear Support Vector Machines

- What if decision boundary is not linear?

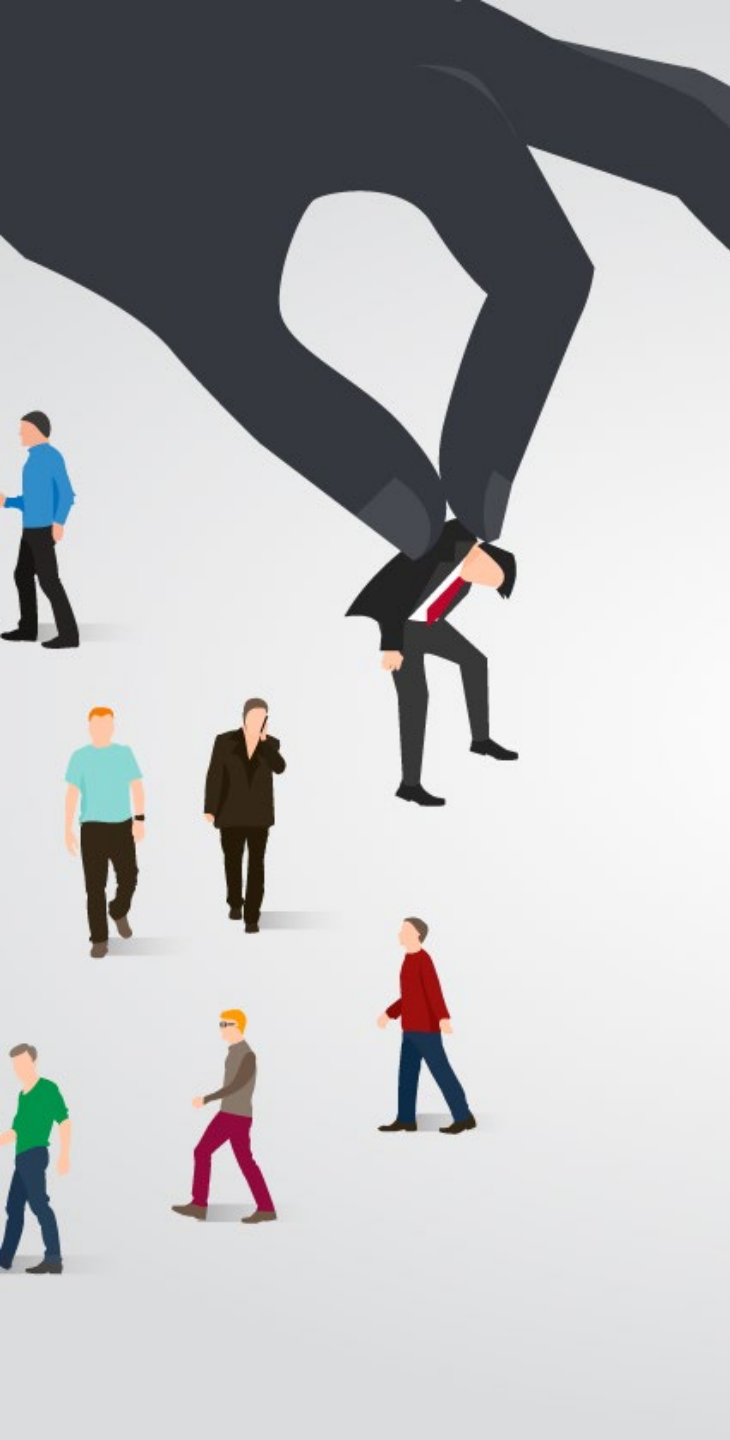




# Nonlinear Support Vector Machines



- Project data into a higher dimensional space where the classes are linearly separable.
- Using the Kernel trick!



# Topics

- Rule-Based Classifier
- Nearest Neighbor Classifier
- Naive Bayes Classifier
- Artificial Neural Networks
- Support Vector Machines
- **Ensemble Methods**

# Ensemble Methods


## Method

1. Construct a set of (possibly weak) classifiers from the training data
2. Predict class label of previously unseen records by aggregating predictions made by multiple classifiers

## Advantages

- Improve the stability and often also the accuracy of classifiers.
- Reduces variance in the prediction
- Reduces overfitting

# General Idea



# Why does it work?

- Suppose there are 25 base classifiers
  - Each classifier has error rate,  $\epsilon = 0.35$
  - Assume classifiers are independent (different features and/or training data)
  - Probability that the ensemble classifier makes a wrong prediction:

$$\sum_{i=13}^{25} \binom{25}{i} \epsilon^i (1 - \epsilon)^{25-i} = 0.06$$

= Probability that 13 or more classifier make the wrong decision

## Notes

- 13 is the majority vote
- The binomial coefficient gives the number of ways you can choose  $i$  out of 25

# Examples of Ensemble Methods

- How to generate an ensemble of classifiers?
  - Bagging
  - Boosting
  - Random Forests

# Bagging (Bootstrap Aggregation)

## 1. **Sampling** with replacement (bootstrap sampling)

<b>Original Data</b>	1	2	3	4	5	6	7	8	9	10
<b>Bagging (Round 1)</b>	7	8	10	8	2	5	10	10	5	9
<b>Bagging (Round 2)</b>	1	4	9	1	2	3	2	7	3	2
<b>Bagging (Round 3)</b>	1	8	5	10	5	5	9	6	3	7

Note: some objects are chosen multiple times in a bootstrap sample while others are not chosen! A typical bootstrap sample contains about 63% of the objects in the original data.

2. **Build classifiers**, one for each bootstrap sample (classifiers are hopefully independent since they are learned from different subsets of the data)

3. **Aggregate** the classifiers' results by averaging or voting

# Boosting


- Records that are incorrectly classified in one round will have their weights increased in the next

Original Data	1	2	3	4	5	6	7	8	9	10
Boosting (Round 1)	7	3	2	8	7	9	4	10	6	3
Boosting (Round 2)	5	4	9	4	2	5	1	7	4	2
Boosting (Round 3)	4	4	8	10	4	5	4	6	3	4

- Example 4 is hard to classify. Its weight is increased; therefore it is more likely to be chosen again in subsequent rounds
- Popular algorithm:** AdaBoost (Adaptive Boosting) typically uses decision trees as the weak learner.




# Random Forests



- Introduce two sources of randomness: “Bagging” and “Random input vectors”
- **Bagging method:** each tree is grown using a bootstrap sample of training data
- **Random vector method:** At each node, best split is chosen only from a random sample of the  $m$  possible attributes.

# Gradient Boosted Decision Trees (XGBoost)




- **Idea:** build models to predict (correct) errors (= boosting).

- **Approach:**

1. Start with a naive (weak) model
2. Calculate errors for each observation in the dataset.
3. Build a new model to predict these errors and add to the ensemble.
4. Go to 2.

# Other Popular Approaches

- Logistic Regression
- Linear Discriminant Analysis
- Regularized Models (Shrinkage)
- Stacking



# Topics

- Other Classification Methods
  - Rule-Based Classifier
  - Nearest Neighbor Classifier
  - Naive Bayes Classifier
  - Artificial Neural Networks
  - Support Vector Machines
  - Ensemble Methods
- **Class Imbalance Problem**

# Class Imbalance Problem

Consider a 2-class problem


- Number of Class 0 examples = 9990
- Number of Class 1 examples = 10

A simple model:

- Always predict Class 0
- accuracy =  $9990/10000 = 99.9\%$
- error =  $0.1\%$

## Issues:

1. Evaluation: accuracy is misleading.
2. Learning: Most classifiers try to optimize accuracy/error. **These classifiers will not learn how to find examples of Class 1!**



# Class Imbalance Problem: Evaluation



**Do not use accuracy to evaluate for problems with strong class imbalance!**

Use instead:

- ROC curves and AUC (area under the curve)
- Precision/Recall plots or the F1 Score
- Cohen's Kappa
- Misclassification cost

# Class Imbalance Problem: Learning

- **Do nothing.** Sometimes you get lucky!
- **Balance the data set:** Down-sample the majority class and/or up-sample the minority class (use sampling with replacement). Synthesize new examples with SMOTE. This will artificially increase the error for a mistake in the minority class.
- Use **algorithms** that can deal with class imbalance (see next slide).
- Throw away minority examples and switch to an **anomaly detection** framework.



# Class Imbalance Problem: Learning

Algorithms that can deal with class imbalance:

- Use a classifier that **predict a probability** and lower the decision threshold (from the default of .5). We can estimate probabilities for decision trees using the positive and negative training examples in each leaf node.
- Use a **cost-sensitive classifier** that considers a cost matrix (not too many are available).
- Use boosting techniques like **AdaBoost**.







## Conclusion

- There are many ways to implement the classification function.
- Each of them has a different inductive bias and often benefits from specifically created feature (e.g., interaction effects in linear models).
- Accuracy is a big problematic for **imbalanced data sets**. Rebalancing the data may be necessary.